

CS445 Final Project

```
In [41]: █ 1 import cv2
2 import numpy as np
3 import os
4 from matplotlib import pyplot as plt
5 from itertools import groupby, combinations
6 %matplotlib inline
```

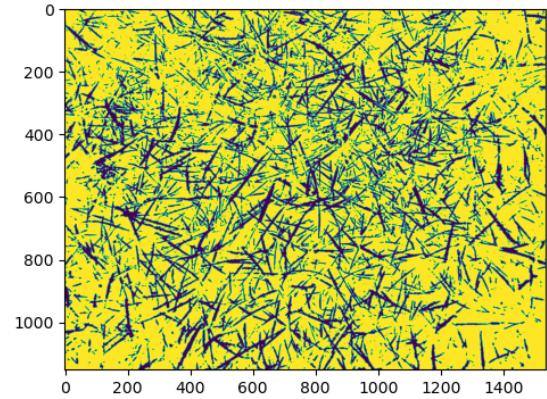
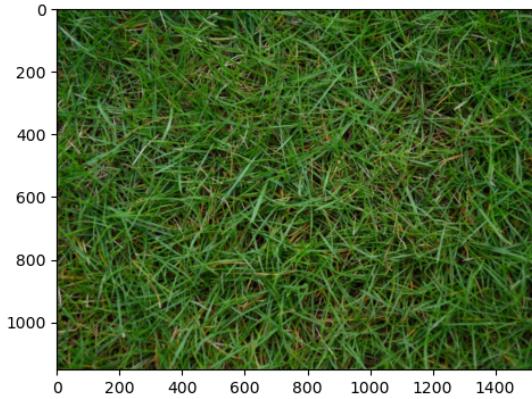
color pixel count

green pixel out of all the pixels. sensitive to color range

```
In [2]: █ 1 def count_grass_pixels(image, lower_green, upper_green):
2     hsv_image = cv2.cvtColor(image, cv2.COLOR_BGR2HSV)
3     mask = cv2.inRange(hsv_image, lower_green, upper_green)
4     count = np.sum(mask > 0)
5     return count, mask
```

```
In [83]: 1 img = cv2.imread('./images/green-bermuda-grass.jpg')
2 lower_green = np.array([40, 0, 0])
3 upper_green = np.array([120, 255, 256])
4
5 hsv_img = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)
6
7 count, mask = count_grass_pixels(hsv_img, lower_green, upper_green)
8 cv2.imwrite("mask.jpg", mask) # save for next parts
9
10 # Calculate the percentage
11 percentage = count / (img.shape[0] * img.shape[1])
12 print(f"score: {percentage*100}%")
13
14 plt.figure(figsize=(12,6))
15
16 plt.subplot(1,2,1)
17 plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
18 plt.plot()
19
20 plt.subplot(1,2,2)
21 plt.imshow(mask)
22 plt.plot();
23
```

score: 75.73281747323496%



lines count (HoughLinesP)

number of line segments to estimate number of blades. Lines have to be very straight. not accurate

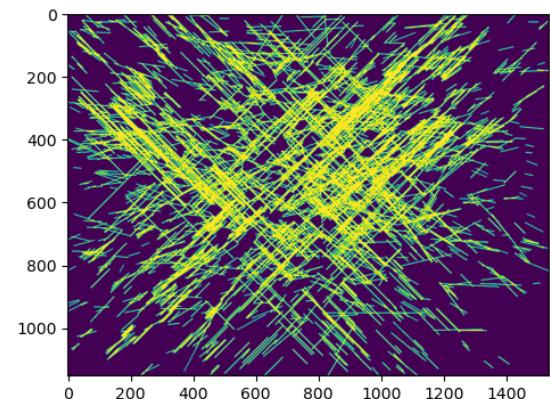
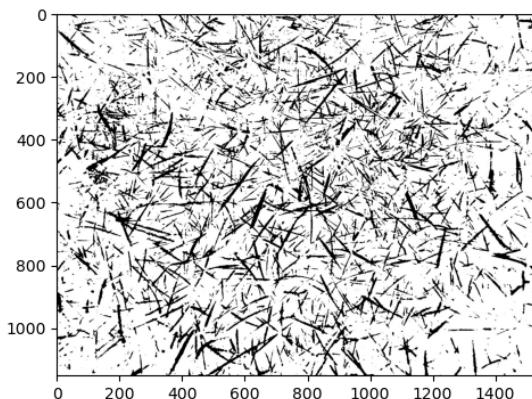
In [4]:

```
1 def count_lines_houghP(img):
2     # Load image
3     mask = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
4
5     # Apply edge detection
6     edges = cv2.Canny(mask, 50, 150)
7
8     # Use Probabilistic Hough Transform to detect line segments
9     lines = cv2.HoughLinesP(edges, 1, np.pi / 180, threshold=150, minL
10
11    # Draw the line segments on the original image
12    lines_img = np.zeros_like(mask)
13    if lines is not None:
14        for line in lines:
15            x1, y1, x2, y2 = line[0]
16            cv2.line(lines_img, (x1, y1), (x2, y2), 255, 2)
17
18    # Count the number of lines
19    num_lines = 0
20    if lines is not None:
21        num_lines = len(lines)
22
23    plt.figure(figsize=(12,6))
24
25    plt.subplot(1,2,1)
26    plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
27    plt.plot()
28
29    plt.subplot(1,2,2)
30    plt.imshow(lines_img)
31    plt.plot();
32
33    print(f'Number of lines: {num_lines}')
34
```

In [5]:

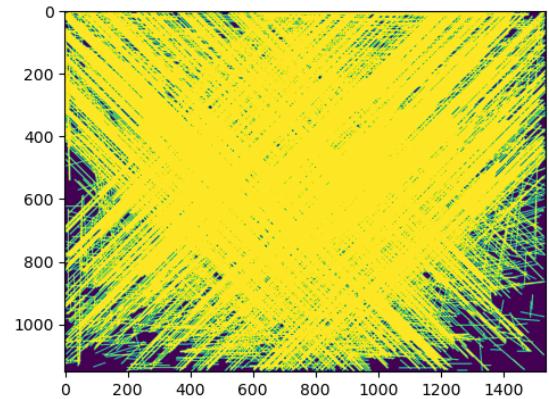
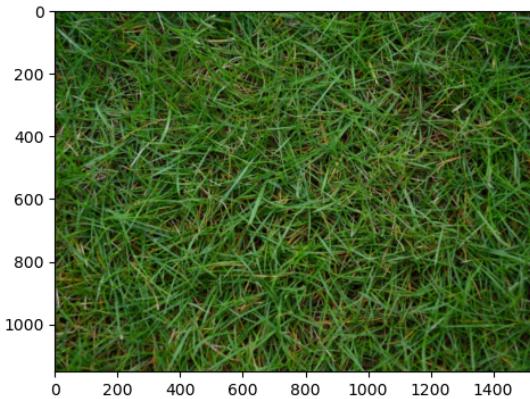
```
1 img = cv2.imread('mask.jpg')
2 count_lines_houghP(img)
3
```

Number of lines: 2225



```
In [6]: 1 img = cv2.imread('./images/green-bermuda-grass.jpg')
2 count_lines_houghP(img)
3
```

Number of lines: 2472



lines count (LineSegmentDetector)

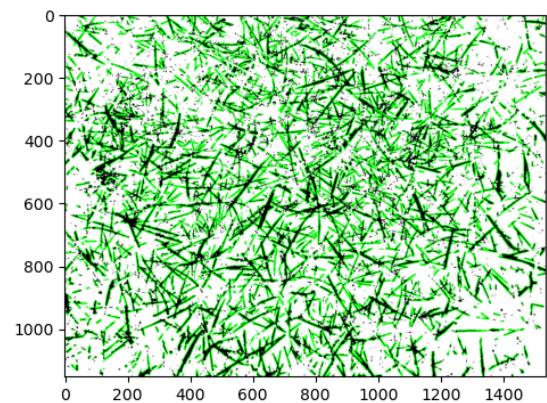
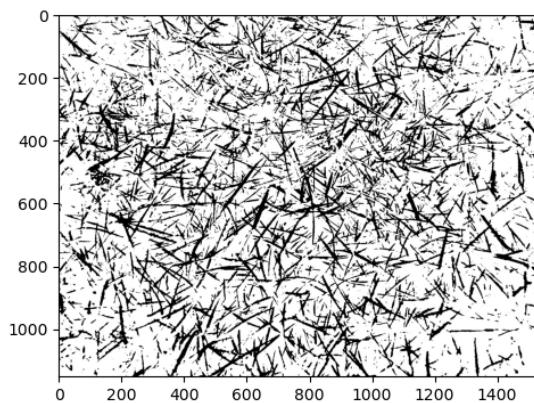
number of line segments to estimate number of blades.

In [192]:

```
1 def count_lines_lsd(img, length_threshold=10):
2     # Load your binary mask image
3     gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
4
5     # Create LineSegmentDetector object
6     lsd = cv2.createLineSegmentDetector()
7     lines, width, prec, nfa = lsd.detect(gray)
8
9     # Draw the line segments on the original image
10    lines_img = np.copy(img)
11    if lines is not None:
12        for line in lines:
13            x1, y1, x2, y2 = map(int, line[0])
14            # Check if the line length is greater than length threshold
15            if np.sqrt((x2 - x1)**2 + (y2 - y1)**2) > length_threshold:
16                cv2.line(lines_img, (x1, y1), (x2, y2), (0, 255, 0), 2)
17
18    # Count the number of lines
19    num_lines = 0
20    if lines is not None:
21        num_lines = sum(1 for line in lines if np.sqrt((line[0][2] - line[0][0])**2 + (line[0][3] - line[0][1])**2) > length_threshold)
22
23    plt.figure(figsize=(12, 6))
24
25    plt.subplot(1, 2, 1)
26    plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
27
28    plt.subplot(1, 2, 2)
29    plt.imshow(cv2.cvtColor(lines_img, cv2.COLOR_BGR2RGB))
30
31    plt.show()
32
33    return round(num_lines / 25) # calibrated from a simple image
34
```

In [178]:

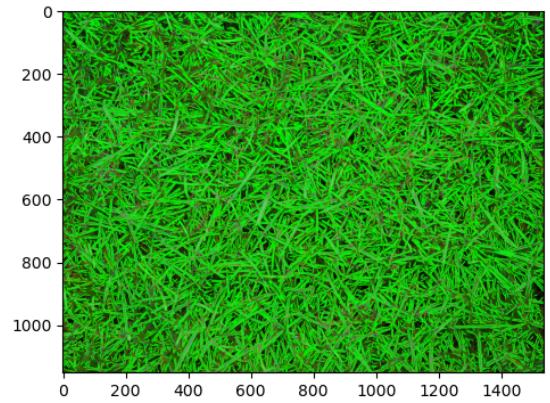
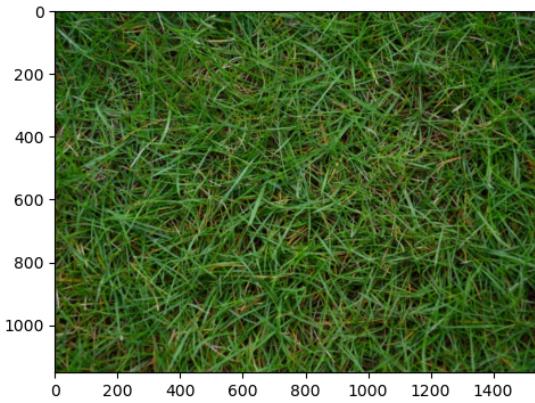
```
1 img = cv2.imread('mask.jpg')
2 print(f'Number of lines: {count_lines_lsd(img)}')
```



Number of lines: 308

In [179]:

```
1 img = cv2.imread('./images/green-bermuda-grass.jpg')
2 print(f'Number of lines: {count_lines_lsd(img)}')
3
```



Number of lines: 571

Calibration and Area

find the area by a known measurement in the image (12 in is ~ 30.48 cm) and dpi

In [180]:

```
1 img = cv2.imread('./images/12in.jpg')
2 img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
3 plt.imshow(img)
4 plt.plot();
5 m_per_pixel = 0.3048 / img.shape[0]
6
7 print(f"every pixel is approx {m_per_pixel} m long, or {(m_per_pixel
8
9 area = img.shape[0] * m_per_pixel * img.shape[1] * m_per_pixel
10 print(f"area of this image is approx {area} m^2")
```

every pixel is approx 7.55952380952381e-05 m long, or 5.714640022675738e-09 m²
area of this image is approx 0.06967728000000001 m²



threshold & adaptive threshold

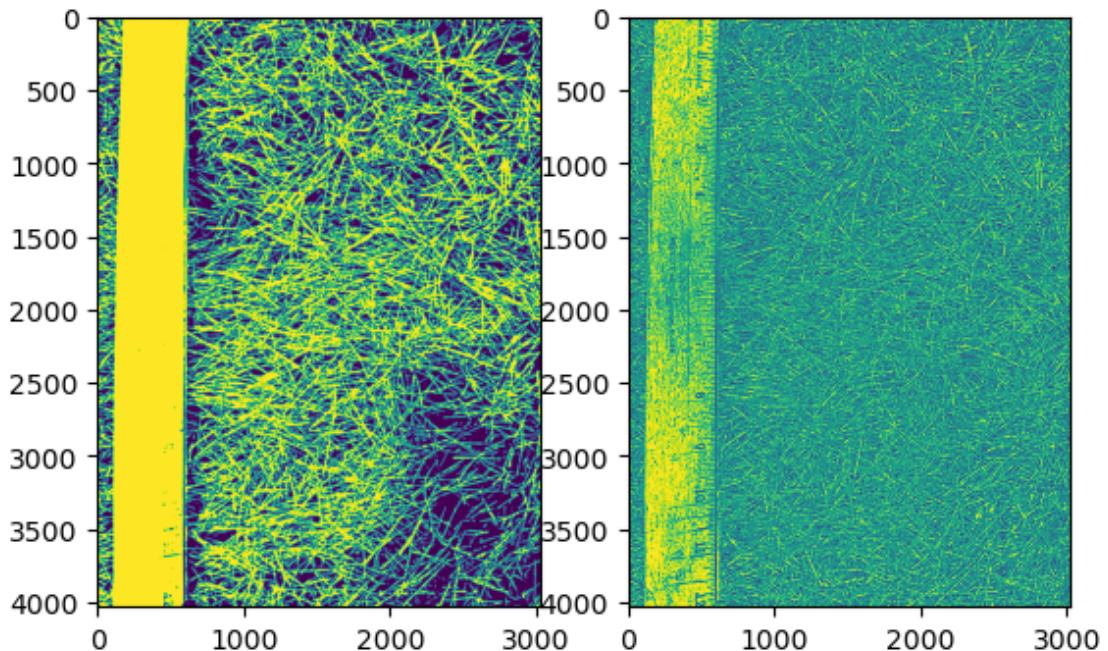
just trying different things out

In [181]:

```
1 def thresholds(img):
2     img = cv2.medianBlur(img,5)
3     _, th1 = cv2.threshold(img,100,255,cv2.THRESH_BINARY)
4     th2 = cv2.adaptiveThreshold(img,255,cv2.ADAPTIVE_THRESH_MEAN_C, cv
5
6     plt.figure()
7
8     plt.subplot(1,2,1)
9     plt.imshow(th1)
10    plt.plot()
11
12    plt.subplot(1,2,2)
13    plt.imshow(th2)
14    plt.plot();
```

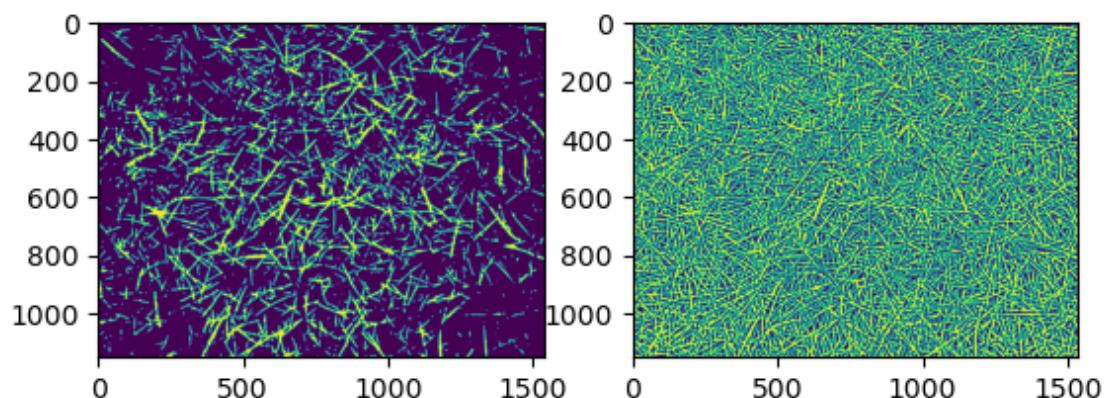
In [182]:

```
1 img = cv2.imread('./images/12in.jpg', cv2.IMREAD_GRAYSCALE)
2 thresholds(img)
```



In [183]:

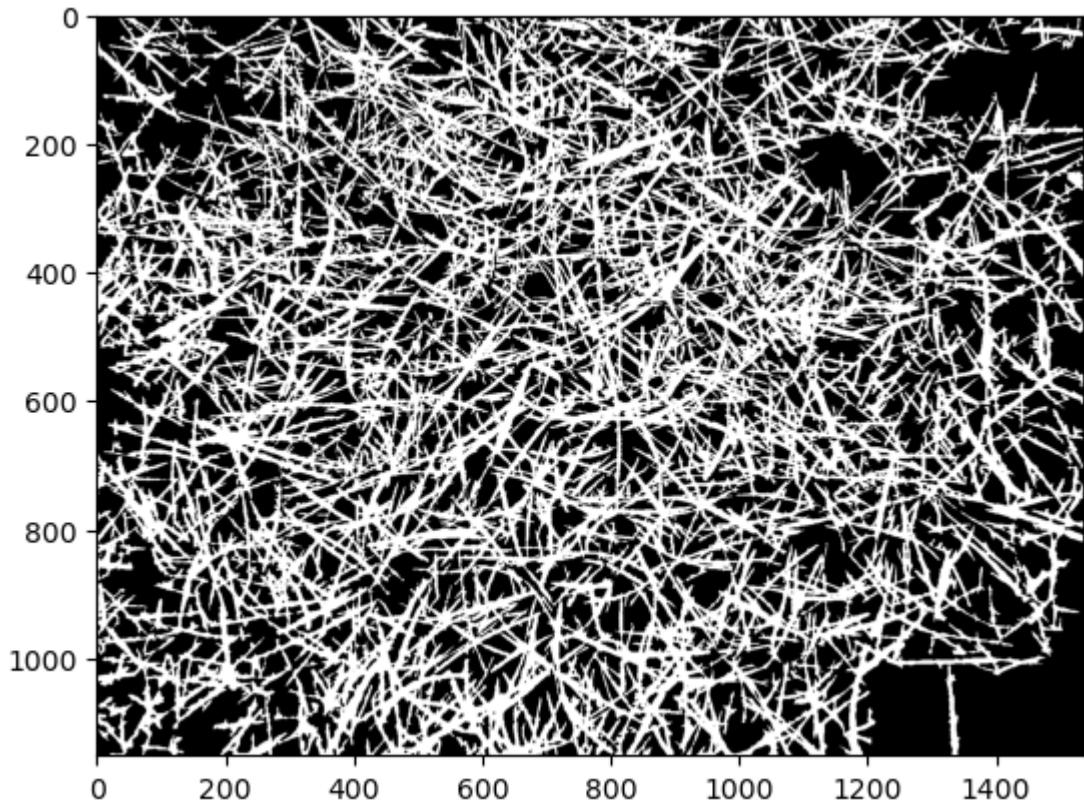
```
1 img = cv2.imread('./images/green-bermuda-grass.jpg', cv2.IMREAD_GRAYSC
2 thresholds(img)
```



In [184]:

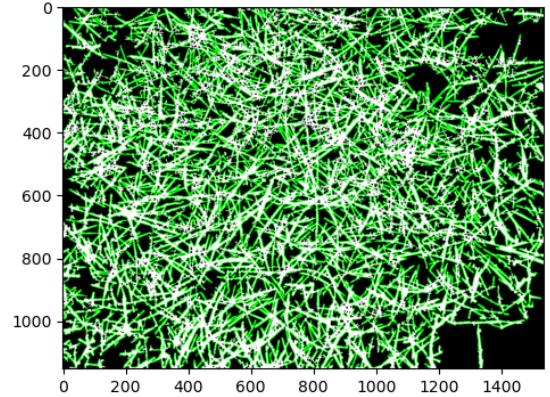
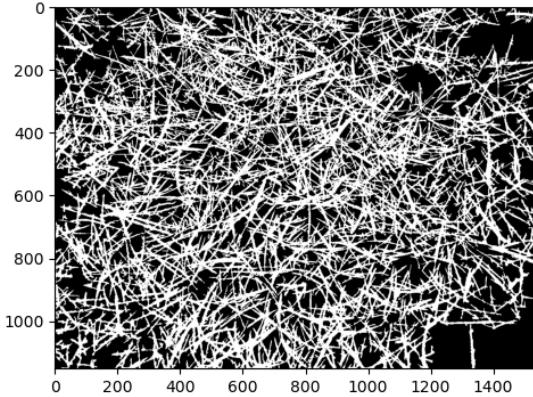
```
1 image = cv2.imread('./images/green-bermuda-grass.jpg', cv2.IMREAD_UNCH
2 gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
3
4 # Use cv2.THRESH_OTSU as a flag to automatically determine the threshold
5 _, binary = cv2.threshold(gray, 0, 255, cv2.THRESH_BINARY | cv2.THRESH_OTSU)
6
7 # getting mask with connectedComponents
8 num_labels, labels, stats, centroids = cv2.connectedComponentsWithStats(
9 mask = np.zeros_like(labels, dtype=np.uint8)
10 mask[labels == 1] = 255
11 print('Area:', stats[1, cv2.CC_STAT_AREA])
12
13 plt.figure()
14 plt.imshow(mask, cmap='gray')
15 plt.show();
16 cv2.imwrite("mask_cc.jpg", mask);
17
```

Area: 780630



In [185]:

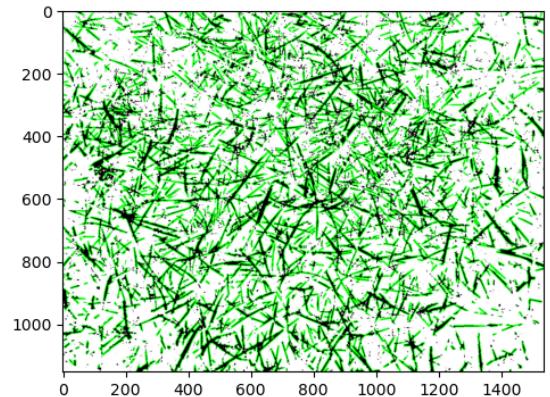
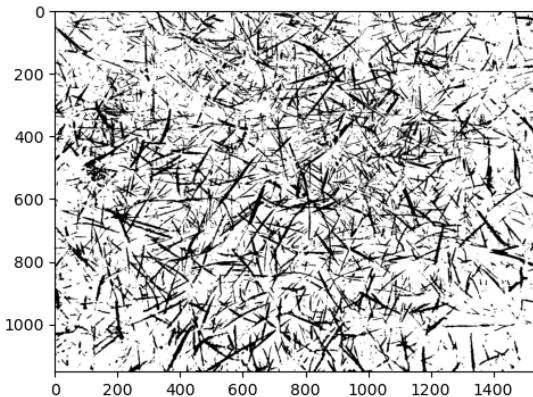
```
1 img = cv2.imread('mask_cc.jpg')
2 lines = count_lines_lsd(img)
3 H, W, _ = img.shape
4 print(f'density is approx {lines} lines / {(10000*H*W*(m_per_pixel**2)):.2f} cm^2')
```



density is approx 398 lines / 101.11895510204084 cm² = 3.94 blades/cm²

In [186]:

```
1 img = cv2.imread('mask.jpg')
2 lines = count_lines_lsd(img)
3 H, W, _ = img.shape
4 print(f'density is approx {lines / (10000*H*W*(m_per_pixel**2)):.2f} blades/cm^2')
```



density is approx 3.05 blades/cm²

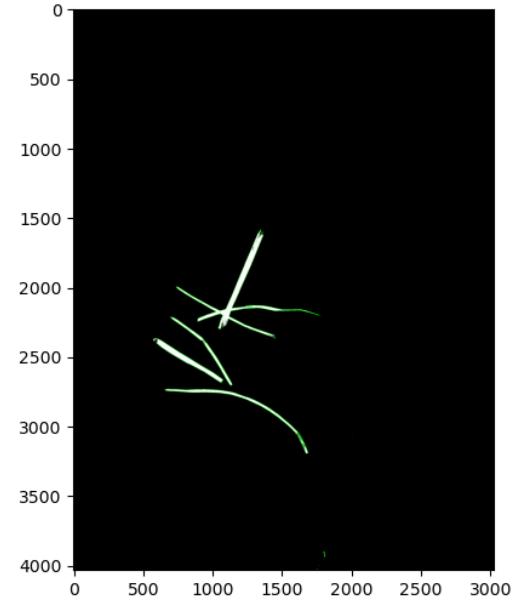
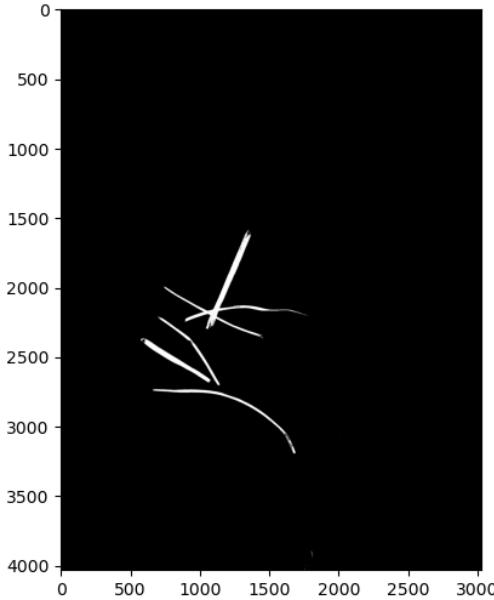
make it a little more organized / streamlined

In [235]:

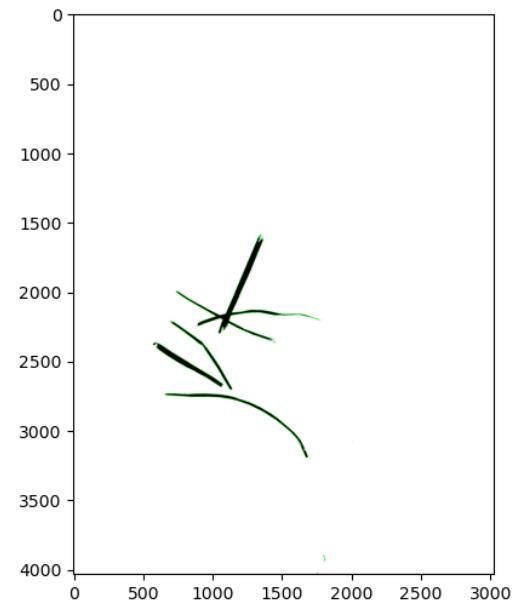
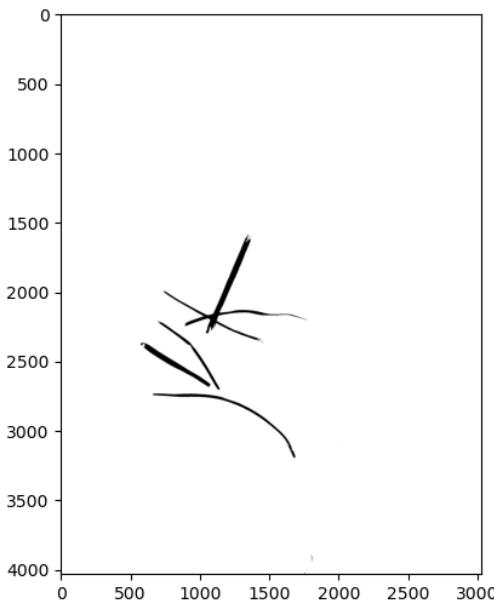
```
1 def density_pipeline(img, area_img, ruler_cm=30.48):
2     # mask from counting green pixels
3     lower_green = np.array([40, 40, 40])
4     upper_green = np.array([120, 255, 256])
5     hsv_img = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)
6     count, mask = count_grass_pixels(hsv_img, lower_green, upper_green)
7     cv2.imwrite('./results/mask.jpg', mask)
8
9     # mask from connected components
10    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
11    _, binary = cv2.threshold(gray, 0, 255, cv2.THRESH_BINARY | cv2.THRESH_OTSU)
12    num_labels, labels, stats, centroids = cv2.connectedComponentsWithStats(binary)
13    max_area_label = np.argmax(stats[1:, cv2.CC_STAT_AREA]) + 1
14    mask_cc = np.zeros_like(labels, dtype=np.uint8)
15    mask_cc[labels == max_area_label] = 255
16    area_cc = stats[max_area_label, cv2.CC_STAT_AREA]
17    cv2.imwrite('./results/mask_cc.jpg', mask_cc)
18
19    # area from ruler
20    cm_per_pixel = ruler_cm / area_img.shape[0]
21    cm2_per_pixel = cm_per_pixel ** 2
22
23    # find lines for mask and mask_cc
24    mask = cv2.imread('./results/mask.jpg')
25    lines = count_lines_lsd(mask)
26    H, W, _ = mask.shape
27    print(f'pixel color: density is approx {lines} lines / {(H*W*cm2_per_pixel)} cm²')
28
29    mask_cc = cv2.imread('./results/mask_cc.jpg')
30    lines = count_lines_lsd(mask_cc)
31    H, W, _ = mask_cc.shape
32    print(f'connected component: density is approx {lines} lines / {(H*W*cm2_per_pixel)} cm²')
```

In [236]:

```
1 img = cv2.imread('./images/bg_overlap_adjusted.jpg')
2 ruler_img = cv2.imread('./images/bg_overlap_12in.jpg')
3
4 assert img is not None, "file could not be read, check with os.path.ex
5 assert ruler_img is not None, "file could not be read, check with os.p
6
7 density_pipeline(img, ruler_img, 24.0)
```



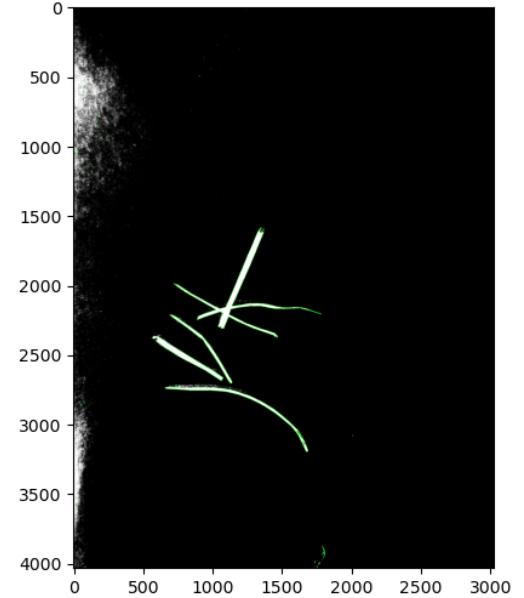
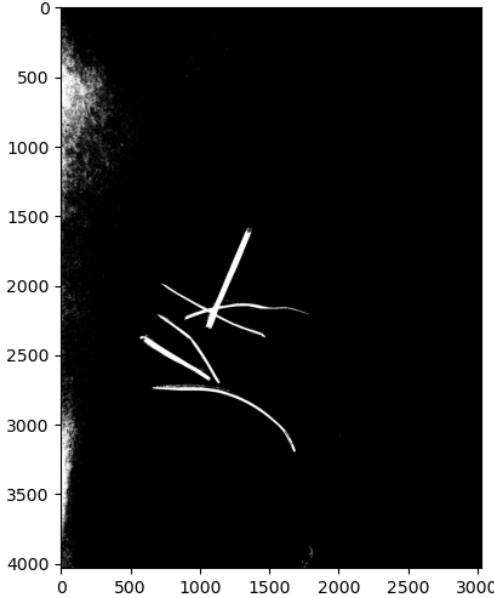
pixel color: density is approx 8 lines / $432.00 \text{ cm}^2 = 0.02 \text{ blades/cm}^2$



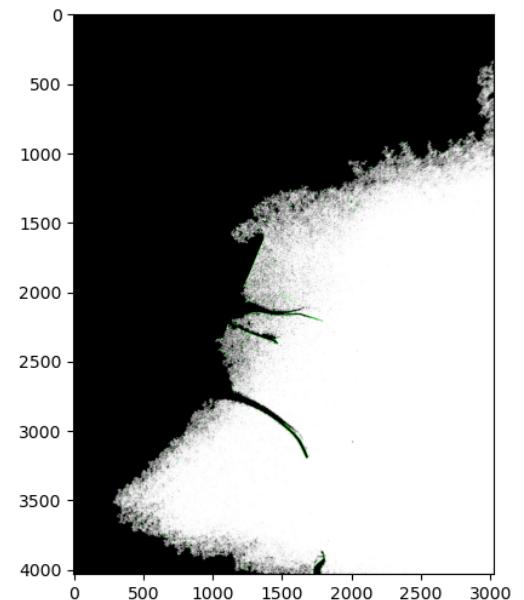
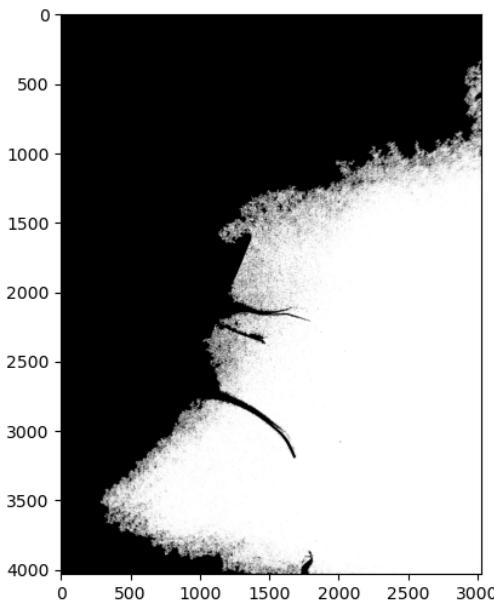
connected component: density is approx 7 lines / $428.37 \text{ cm}^2 = 0.02 \text{ blad}es/\text{cm}^2$

In [237]:

```
1 img = cv2.imread('./images/bg_overlap.jpg')
2 ruler_img = cv2.imread('./images/bg_overlap_12in.jpg')
3
4 assert img is not None, "file could not be read, check with os.path.ex
5 assert ruler_img is not None, "file could not be read, check with os.p
6
7 density_pipeline(img, ruler_img)
```



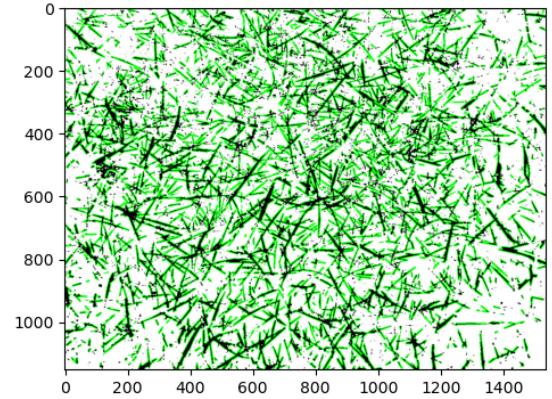
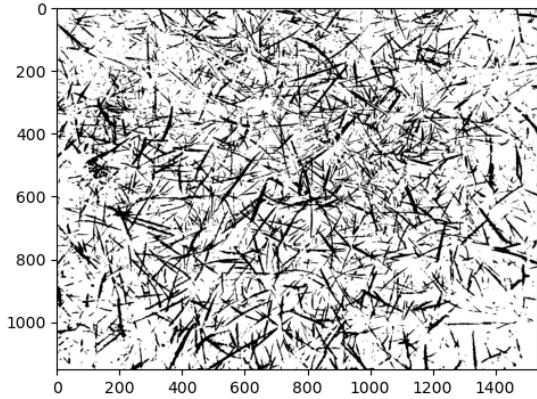
pixel color: density is approx 11 lines / $696.77 \text{ cm}^2 = 0.02 \text{ blades/cm}^2$



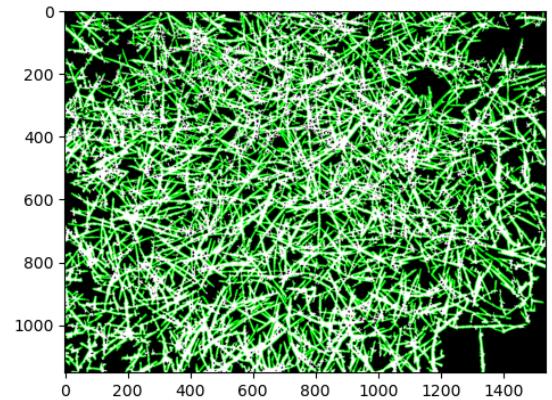
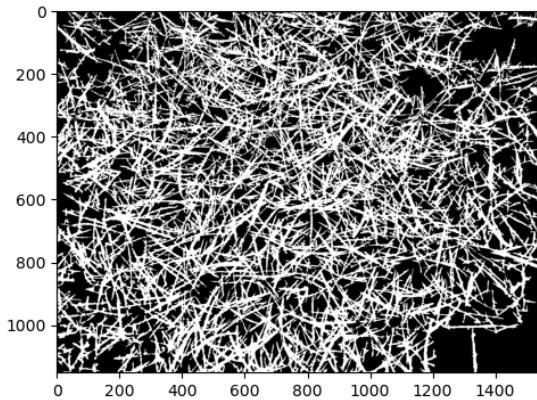
connected component: density is approx 17 lines / $320.45 \text{ cm}^2 = 0.05 \text{ b}$
lades/cm²

In [238]:

```
1 img = cv2.imread('./images/green-bermuda-grass.jpg')
2 ruler_img = cv2.imread('./images/12in.jpg')
3
4 assert img is not None, "file could not be read, check with os.path.ex
5 assert ruler_img is not None, "file could not be read, check with os.p
6
7 density_pipeline(img, ruler_img)
```



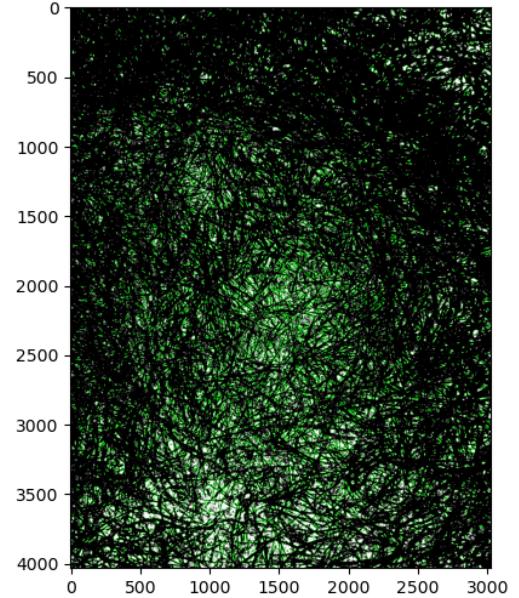
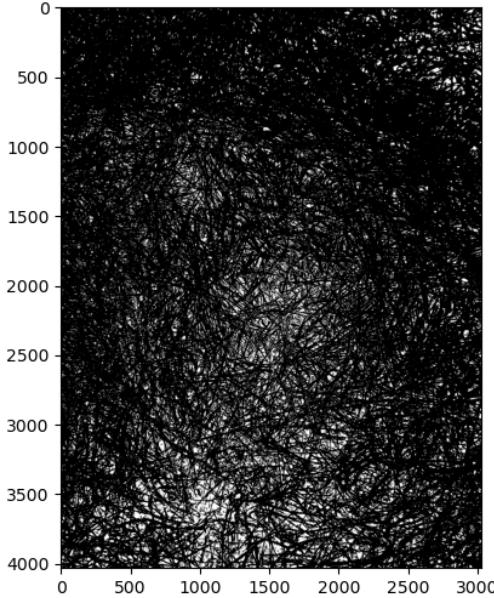
pixel color: density is approx 308 lines / $101.12 \text{ cm}^2 = 3.05 \text{ blades/cm}^2$



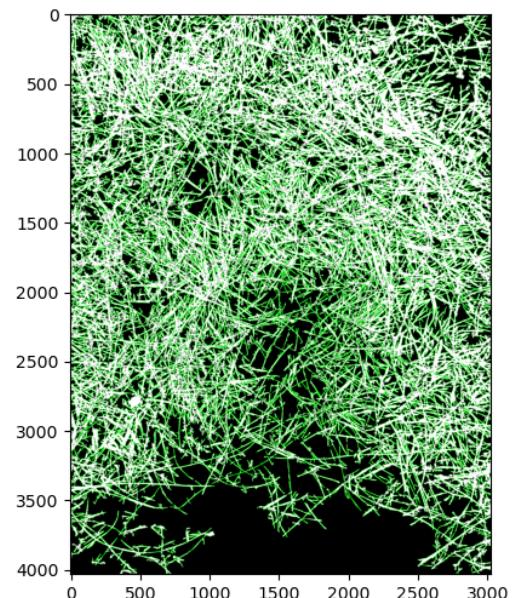
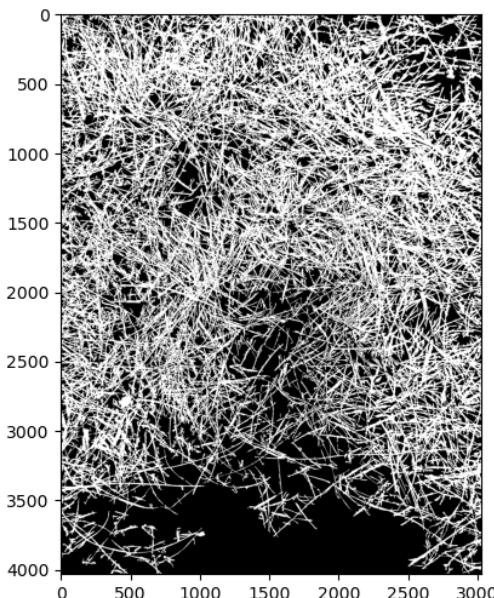
connected component: density is approx 398 lines / $44.61 \text{ cm}^2 = 8.92 \text{ b}$ lades/cm²

In [239]:

```
1 img = cv2.imread('./images/inside.jpg')
2 ruler_img = cv2.imread('./images/inside_12in.jpg')
3
4 assert img is not None, "file could not be read, check with os.path.exists"
5 assert ruler_img is not None, "file could not be read, check with os.path.exists"
6
7 density_pipeline(img, ruler_img)
```



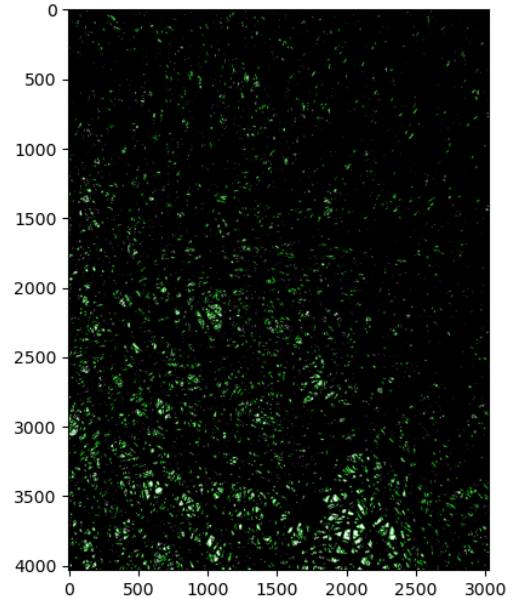
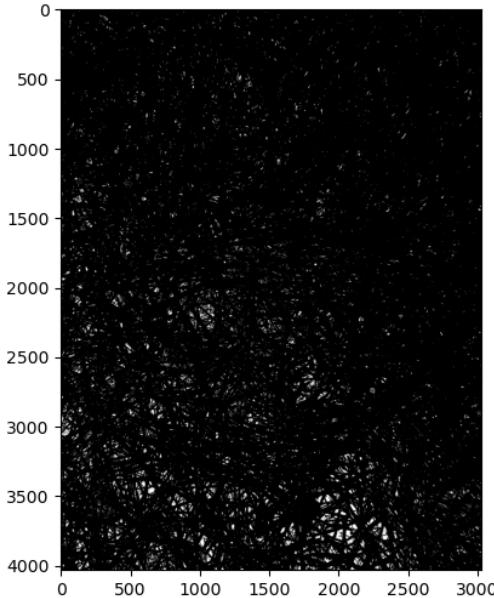
pixel color: density is approx 731 lines / 696.77 cm^2 = 1.05 blades/cm 2



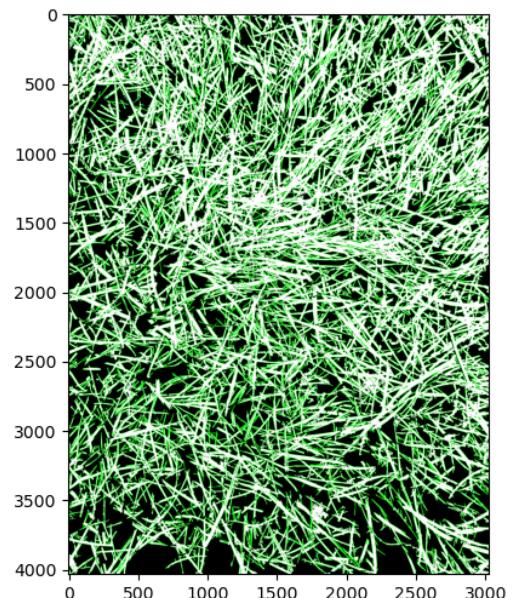
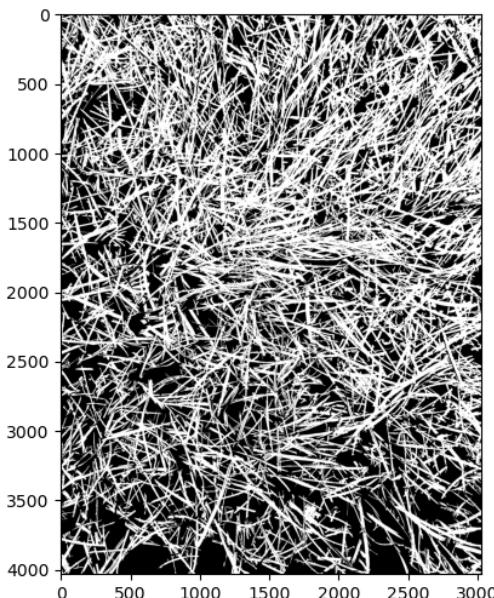
connected component: density is approx 1591 lines / 337.80 cm^2 = 4.71 blades/cm 2

In [240]:

```
1 img = cv2.imread('./images/outside.jpg')
2 ruler_img = cv2.imread('./images/outside_12in.jpg')
3
4 assert img is not None, "file could not be read, check with os.path.exists"
5 assert ruler_img is not None, "file could not be read, check with os.path.exists"
6
7 density_pipeline(img, ruler_img)
```



pixel color: density is approx 191 lines / $696.77 \text{ cm}^2 = 0.27 \text{ blades/cm}^2$



connected component: density is approx 1303 lines / $352.35 \text{ cm}^2 = 3.70 \text{ blades/cm}^2$

lines counted using LSD does not equal to # of blades, but limiting the length of lines helped estimate the density calculation

which approach gives results closer to actual value or human guesses?

come back to this later: resolving the overlaps

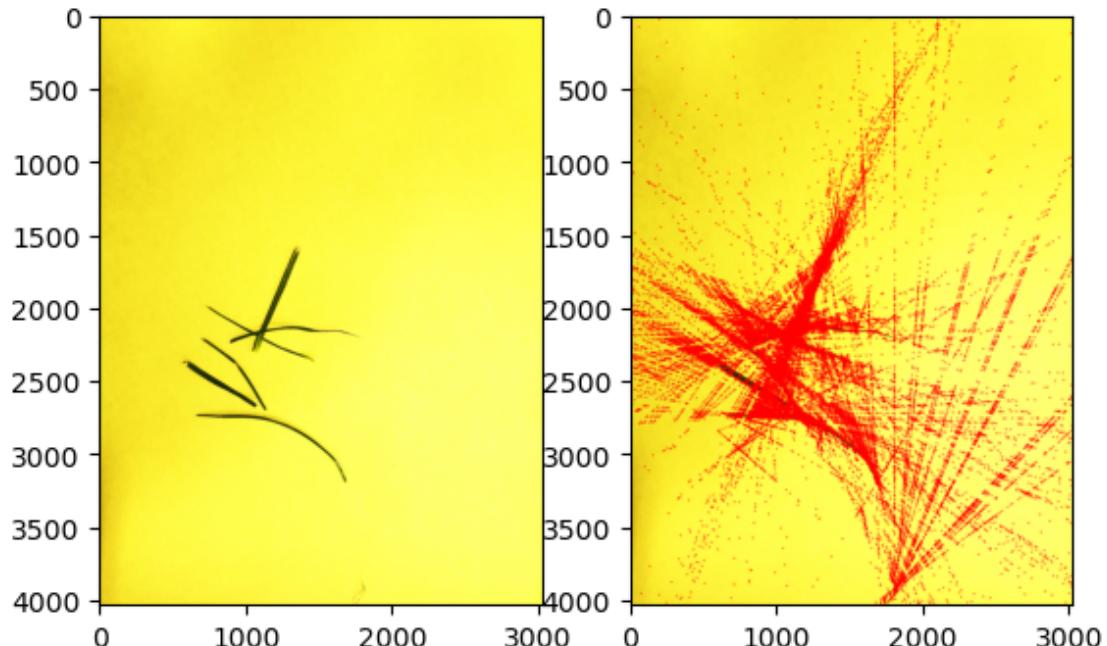
since lsd counts a short line, how to count a curvy grass blade as 1?

```
In [78]: ┌─ 1 def find_line_intersections(lines):
  2     intersections = []
  3     for line1, line2 in combinations(lines, 2):
  4         x1, y1, x2, y2 = line1[0]
  5         x3, y3, x4, y4 = line2[0]
  6
  7         # Check if Lines are not parallel
  8         det = (x1 - x2) * (y3 - y4) - (y1 - y2) * (x3 - x4)
  9
 10        if det != 0:
 11            intersection_x = ((x1 * y2 - y1 * x2) * (x3 - x4) - (x1 -
 12            intersection_y = ((x1 * y2 - y1 * x2) * (y3 - y4) - (y1 -
 13
 14            intersections.append((int(intersection_x), int(intersection_y)))
 15
 16    return intersections
 17
```

In [139]:

```
1 img = cv2.imread('./images/bg_overlap_adjusted.jpg')
2 gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
3
4 # Use the LineSegmentDetector to detect lines
5 lsd = cv2.createLineSegmentDetector()
6 lines, _, _, _ = lsd.detect(gray)
7
8 # Draw the detected lines on the image
9 lines_img = np.copy(img)
10 for line in lines:
11     x1, y1, x2, y2 = map(int, line[0])
12     cv2.line(lines_img, (x1, y1), (x2, y2), (0, 255, 0), 2)
13
14 intersections = find_line_intersections(lines)
15
16 # Draw the intersections on the image
17 for point in intersections:
18     cv2.circle(lines_img, point, 5, (0, 0, 255), 2)
19
20 print(len(intersections), len(lines))
21 plt.figure()
22
23 plt.subplot(1,2,1)
24 plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
25 plt.plot()
26
27 plt.subplot(1,2,2)
28 plt.imshow(cv2.cvtColor(lines_img, cv2.COLOR_BGR2RGB))
29 plt.plot();
30
```

82621 407



come back to this later: compare to known density patches

In []: 1

come back to this later: use projection from a video like proj5

In []: 1

come back to this later: TP-LSD

<https://github.com/Siyuada7/TP-LSD> (<https://github.com/Siyuada7/TP-LSD>)

In []: 1

Graveyard

Algorithm from the paper: Run-Length Encoding (RLE)

searching for contiguous pixels along the major axis in an input binary image and replacing them by a single instance of the run count for each set of contiguous pixels until the end of the image is reached.

- Step 1: read image and remove artifacts and extraneous material that do not meet the minimum threshold pixel requirement
- Step 2: Set minimum and maximum thresholds for the contiguous pixel widths
- Step 3: Create a zero matrix of the image size with a one pixel buffer
- Step 4: Assign original matrix to the zero matrix
- Step 5: Create a vector matrix from the transpose of the original matrix
- Step 6: Find positions where one element is different from the next
- Step 7: Find run lengths and values after determining the end position of each run
- Step 8: Assign zeros for all widths above and below the specified threshold range
- Step 9: Perform run-length decoding
- Step 10: Repeat Steps 5 through 8 for each image angle from 0 to 90 degrees at 5 degrees interval
- Step 11: Calculate image statistics

In [28]:

```
1 def preprocess_image(img, threshold=128):
2     # Step 1: Read image and remove artifacts
3     _, binary_image = cv2.threshold(img, threshold, 255, cv2.THRESH_BINARY)
4     return binary_image
5
6 def run_length_encoding(binary_image, min_threshold=5, max_threshold=1000):
7     # Step 3: Create a zero matrix with a one-pixel buffer
8     zero_matrix = np.zeros((binary_image.shape[0] + 2, binary_image.shape[1] + 2))
9
10    # Step 4: Assign original matrix to the zero matrix
11    zero_matrix[1:-1, 1:-1] = binary_image
12
13    encoded_runs = []
14
15    # Step 10: Repeat Steps 5 through 8 for each image angle from 0 to 90 degrees
16    for angle in range(0, 91, 5):
17        rotated_image = rotate_image(binary_image, angle)
18        runs = find_runs(rotated_image)
19        filtered_runs = filter_runs(runs, min_threshold, max_threshold)
20        decoded_image = run_length_decoding(filtered_runs)
21
22        # Step 11: Calculate image statistics
23        calculate_image_statistics(decoded_image)
24        print(decoded_image)
25
26        plt.figure()
27        plt.imshow(decoded_image)
28        plt.show();
29
30 def rotate_image(image, angle):
31     # Step 5: Create a vector matrix from the transpose of the original image
32     rotated_image = np.transpose(image)
33
34     # Rotate the image
35     rotation_matrix = cv2.getRotationMatrix2D((rotated_image.shape[1] / 2, rotated_image.shape[0] / 2), angle, 1)
36     rotated_image = cv2.warpAffine(rotated_image, rotation_matrix, (rotated_image.shape[1], rotated_image.shape[0]))
37
38     return rotated_image
39
40 def find_runs(image):
41     # Step 6: Find positions where one element is different from the next
42     runs = []
43     for row in image:
44         run_lengths = [sum(1 for _ in group) for key, group in groupby(row)]
45         run_values = [key for key, group in groupby(row)]
46         runs.append((run_lengths, run_values))
47
48     return runs
49
50 def filter_runs(runs, min_threshold, max_threshold):
51     # Step 7: Find run lengths and values after determining the end points
52     filtered_runs = []
53     for run_lengths, run_values in runs:
54         filtered_lengths = [length if min_threshold <= length <= max_threshold else 0 for length in run_lengths]
55         filtered_runs.append((filtered_lengths, run_values))
```

```

56
57     return filtered_runs
58
59 def run_length_decoding(runs):
60     # Step 8: Assign zeros for all widths above and below the specified
61     decoded_image = np.zeros(len(runs), sum(sum(run_lengths) for run_
62
63     # Perform run-Length decoding
64     for i, (run_lengths, run_values) in enumerate(runs):
65         start = 0
66         for length, value in zip(run_lengths, run_values):
67             decoded_image[i, start:start + length] = value
68             start += length
69
70     return decoded_image
71
72 def calculate_image_statistics(image):
73     # Implement image statistics calculation (mean, median, etc.) here
74     pass

```

In [31]:

```

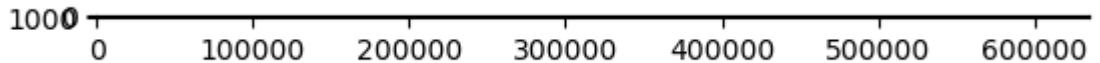
1 # Specify the path to your image
2 image = cv2.imread('./images/green-bermuda-grass.jpg', cv2.IMREAD_GRAYSCALE)
3
4 # Step 2: Set minimum and maximum thresholds for the contiguous pixel
5 min_threshold = 5
6 max_threshold = 100
7
8 binary_image = preprocess_image(image)
9 run_length_encoding(binary_image, min_threshold, max_threshold)
10

```

```

[[255 255 255 ... 0 0 0]
 [255 255 255 ... 0 0 0]
 [255 255 255 ... 0 0 0]
 ...
 [ 0  0  0 ... 0 0 0]
 [ 0  0  0 ... 0 0 0]
 [ 0  0  0 ... 0 0 0]]

```



```

[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 ...
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]]

```

In []:

1